

binary-classification-xgboost-1

November 15, 2024

0.1 Preprocessing the Train Data

```
[29]: # This Python 3 environment comes with many helpful analytics libraries
      ↪ installed
      # It is defined by the kaggle/python Docker image: https://github.com/kaggle/
      ↪ docker-python
      # For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list
      ↪ all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that
      ↪ gets preserved as output when you create a version using "Save & Run All"
# You can also write temporary files to /kaggle/temp/, but they won't be saved
      ↪ outside of the current session
```

```
/kaggle/input/playground-series-s4e7/sample_submission.csv
/kaggle/input/playground-series-s4e7/train.csv
/kaggle/input/playground-series-s4e7/test.csv
```

```
[30]: df = pd.read_csv("/kaggle/input/playground-series-s4e7/train.csv")
      df.head()
```

```
[30]:   id  Gender  Age  Driving_License  Region_Code  Previously_Insured  \
0    0   Male   21                1           35.0                0
1    1   Male   43                1           28.0                0
2    2  Female   25                1           14.0                1
3    3  Female   35                1            1.0                0
4    4  Female   36                1           15.0                1
```

	Vehicle_Age	Vehicle_Damage	Annual_Premium	Policy_Sales_Channel	Vintage	\
0	1-2 Year	Yes	65101.0	124.0	187	
1	> 2 Years	Yes	58911.0	26.0	288	
2	< 1 Year	No	38043.0	152.0	254	
3	1-2 Year	Yes	2630.0	156.0	76	
4	1-2 Year	No	31951.0	152.0	294	

	Response
0	0
1	1
2	0
3	0
4	0

0.1.1 Checking the null values in the columns

```
[31]: df.isnull().sum()
```

```
[31]: id                0
      Gender            0
      Age              0
      Driving_License   0
      Region_Code       0
      Previously_Insured 0
      Vehicle_Age       0
      Vehicle_Damage    0
      Annual_Premium    0
      Policy_Sales_Channel 0
      Vintage           0
      Response          0
      dtype: int64
```

0.1.2 Unique Value count

```
[32]: df.nunique()
```

```
[32]: id                11504798
      Gender              2
      Age                66
      Driving_License     2
      Region_Code        54
      Previously_Insured   2
      Vehicle_Age         3
      Vehicle_Damage      2
      Annual_Premium      51728
```

```
Policy_Sales_Channel      152
Vintage                  290
Response                   2
dtype: int64
```

0.1.3 Description of the Dataset

```
[33]: df.describe()
```

```
[33]:
```

	id	Age	Driving_License	Region_Code	\
count	1.150480e+07	1.150480e+07	1.150480e+07	1.150480e+07	
mean	5.752398e+06	3.838356e+01	9.980220e-01	2.641869e+01	
std	3.321149e+06	1.499346e+01	4.443120e-02	1.299159e+01	
min	0.000000e+00	2.000000e+01	0.000000e+00	0.000000e+00	
25%	2.876199e+06	2.400000e+01	1.000000e+00	1.500000e+01	
50%	5.752398e+06	3.600000e+01	1.000000e+00	2.800000e+01	
75%	8.628598e+06	4.900000e+01	1.000000e+00	3.500000e+01	
max	1.150480e+07	8.500000e+01	1.000000e+00	5.200000e+01	

	Previously_Insured	Annual_Premium	Policy_Sales_Channel	Vintage	\
count	1.150480e+07	1.150480e+07	1.150480e+07	1.150480e+07	
mean	4.629966e-01	3.046137e+04	1.124254e+02	1.638977e+02	
std	4.986289e-01	1.645475e+04	5.403571e+01	7.997953e+01	
min	0.000000e+00	2.630000e+03	1.000000e+00	1.000000e+01	
25%	0.000000e+00	2.527700e+04	2.900000e+01	9.900000e+01	
50%	0.000000e+00	3.182400e+04	1.510000e+02	1.660000e+02	
75%	1.000000e+00	3.945100e+04	1.520000e+02	2.320000e+02	
max	1.000000e+00	5.401650e+05	1.630000e+02	2.990000e+02	

	Response
count	1.150480e+07
mean	1.229973e-01
std	3.284341e-01
min	0.000000e+00
25%	0.000000e+00
50%	0.000000e+00
75%	0.000000e+00
max	1.000000e+00

```
[34]: columns_to_drop = [ 'Region_Code' ]

# Using the `drop` method
df_new = df.drop(columns=columns_to_drop)
```

```
[35]: df_new
```

```
[35]:
```

	id	Gender	Age	Driving_License	Previously_Insured	\
0	0	Male	21	1	0	
1	1	Male	43	1	0	
2	2	Female	25	1	1	
3	3	Female	35	1	0	
4	4	Female	36	1	1	
...	
11504793	11504793	Male	48	1	0	
11504794	11504794	Female	26	1	0	
11504795	11504795	Female	29	1	1	
11504796	11504796	Female	51	1	0	
11504797	11504797	Male	25	1	1	

	Vehicle_Age	Vehicle_Damage	Annual_Premium	Policy_Sales_Channel	\
0	1-2 Year	Yes	65101.0	124.0	
1	> 2 Years	Yes	58911.0	26.0	
2	< 1 Year	No	38043.0	152.0	
3	1-2 Year	Yes	2630.0	156.0	
4	1-2 Year	No	31951.0	152.0	
...	
11504793	1-2 Year	Yes	27412.0	26.0	
11504794	< 1 Year	Yes	29509.0	152.0	
11504795	< 1 Year	No	2630.0	152.0	
11504796	1-2 Year	Yes	48443.0	26.0	
11504797	< 1 Year	No	32855.0	152.0	

	Vintage	Response
0	187	0
1	288	1
2	254	0
3	76	0
4	294	0
...
11504793	218	0
11504794	115	1
11504795	189	0
11504796	274	1
11504797	189	0


```
[11504798 rows x 11 columns]
```

0.2 Binning Age and Annual Premium

**** To bin the Age and Annual_Premium columns in the dataset, we first define the respective bins. For Age, we set the bins as [10, 20, 30, 40, 50, 60, 70, 80, 90], and for Annual_Premium, the bins are [0, 10000, 20000, 30000, 40000, 50000, 60000, 70000, 80000, 90000, 100000]. Using the pd.cut function, we categorize the Age data into these specified bins with labels representing the age

ranges, and similarly, categorize the Annual_Premium data into the premium bins with numerical labels. After binning the Age column, we drop the original Age column from the DataFrame to avoid redundancy. The same binning process is applied to the Annual_Premium column. This procedure ensures that the age and premium data are properly categorized into meaningful intervals, facilitating subsequent analysis or modeling.****

```
[36]: age_bins = [10, 20, 30, 40, 50, 60, 70, 80, 90]

# Apply binning using cut
df_new['Age_Binned'] = pd.cut(df_new['Age'], bins=age_bins,
    ↳labels=[f'{age}-{age+10}' for age in age_bins[:-1]])

df_new.drop(columns=['Age'], inplace=True)
```

```
[37]: bins = [0, 10000, 20000, 30000, 40000, 50000, 60000, 70000, 80000, 90000,
    ↳100000]
labels = range(len(bins)-1)

df_new['Annual_Premium_Binned'] = pd.cut(df_new['Annual_Premium'], bins=bins,
    ↳labels=labels)
```

```
[38]: df_new
```

```
[38]:
```

	id	Gender	Driving_License	Previously_Insured	Vehicle_Age	\
0	0	Male	1	0	1-2 Year	
1	1	Male	1	0	> 2 Years	
2	2	Female	1	1	< 1 Year	
3	3	Female	1	0	1-2 Year	
4	4	Female	1	1	1-2 Year	
...	
11504793	11504793	Male	1	0	1-2 Year	
11504794	11504794	Female	1	0	< 1 Year	
11504795	11504795	Female	1	1	< 1 Year	
11504796	11504796	Female	1	0	1-2 Year	
11504797	11504797	Male	1	1	< 1 Year	
	Vehicle_Damage	Annual_Premium	Policy_Sales_Channel	Vintage	\	
0	Yes	65101.0	124.0	187		
1	Yes	58911.0	26.0	288		
2	No	38043.0	152.0	254		
3	Yes	2630.0	156.0	76		
4	No	31951.0	152.0	294		
...		
11504793	Yes	27412.0	26.0	218		
11504794	Yes	29509.0	152.0	115		
11504795	No	2630.0	152.0	189		

11504796	Yes	48443.0	26.0	274
11504797	No	32855.0	152.0	189

	Response	Age_Binned	Annual_Premium_Binned
0	0	20-30	6
1	1	40-50	5
2	0	20-30	3
3	0	30-40	0
4	0	30-40	3
...
11504793	0	40-50	2
11504794	1	20-30	2
11504795	0	20-30	0
11504796	1	50-60	4
11504797	0	20-30	3

[11504798 rows x 12 columns]

```
[39]: df_new.isnull().sum()
```

```
[39]: id                0
      Gender            0
      Driving_License   0
      Previously_Insured 0
      Vehicle_Age       0
      Vehicle_Damage    0
      Annual_Premium     0
      Policy_Sales_Channel 0
      Vintage            0
      Response           0
      Age_Binned         0
      Annual_Premium_Binned 2307
      dtype: int64
```

```
[40]: rand_samp = df_new.sample(n=10)
```

```
rand_samp
```

```
[40]:
```

	id	Gender	Driving_License	Previously_Insured	Vehicle_Age	\
8174007	8174007	Female	1	1	< 1 Year	
2246841	2246841	Male	1	1	1-2 Year	
10180753	10180753	Male	1	0	1-2 Year	
7944036	7944036	Male	1	0	< 1 Year	
6829332	6829332	Female	1	0	1-2 Year	
2770997	2770997	Male	1	1	1-2 Year	
4141434	4141434	Female	1	1	1-2 Year	
6541288	6541288	Female	1	0	1-2 Year	

7021528	7021528	Female	1	1	1-2 Year
601478	601478	Female	1	0	1-2 Year

	Vehicle_Damage	Annual_Premium	Policy_Sales_Channel	Vintage	\
8174007	No	23595.0	152.0	213	
2246841	No	28578.0	124.0	76	
10180753	Yes	2630.0	157.0	299	
7944036	Yes	20229.0	124.0	257	
6829332	Yes	31875.0	124.0	241	
2770997	No	2630.0	26.0	175	
4141434	No	24002.0	152.0	185	
6541288	Yes	29364.0	124.0	223	
7021528	No	52639.0	26.0	187	
601478	Yes	29026.0	124.0	182	

	Response	Age_Binned	Annual_Premium_Binned
8174007	0	30-40	2
2246841	0	40-50	2
10180753	1	30-40	0
7944036	0	20-30	2
6829332	1	30-40	3
2770997	0	30-40	0
4141434	0	30-40	2
6541288	1	30-40	2
7021528	0	50-60	5
601478	0	30-40	2

0.3 Label Encoding

- **Import Libraries:**
 - We import pandas for data manipulation and LabelEncoder from sklearn.preprocessing for label encoding.
- **Define LabelEncoder:**
 - An instance of LabelEncoder is created to perform the encoding.
- **Columns to Encode:**
 - We specify the list `columns_to_encode` containing names of categorical columns ('Gender', 'Vehicle_Age', 'Vehicle_Damage', 'Age_Binned') in `df_new` that we want to encode.
- **Apply Label Encoding:**
 - Using a loop, we iterate through each column in `columns_to_encode` and apply `fit_transform` method of LabelEncoder to convert each categorical column into numerical labels.

```
[41]: import pandas as pd
from sklearn.preprocessing import LabelEncoder

# Assuming df is your DataFrame
```

```

label_encoder = LabelEncoder()

# Columns to encode
columns_to_encode = ['Gender', 'Vehicle_Age', 'Vehicle_Damage', 'Age_Binned']

# Apply label encoding to each column
for column in columns_to_encode:
    df_new[column] = label_encoder.fit_transform(df_new[column])

```

```

[42]: df_new.drop(columns=['Annual_Premium'], inplace=True)
df_new

```

```

[42]:
      id  Gender  Driving_License  Previously_Insured  Vehicle_Age  \
0      0      1      1      0      0
1      1      1      1      0      2
2      2      0      1      1      1
3      3      0      1      0      0
4      4      0      1      1      0
...    ...    ...    ...    ...    ...
11504793  11504793      1      1      0      0
11504794  11504794      0      1      0      1
11504795  11504795      0      1      1      1
11504796  11504796      0      1      0      0
11504797  11504797      1      1      1      1

      Vehicle_Damage  Policy_Sales_Channel  Vintage  Response  Age_Binned  \
0      1      124.0      187      0      1
1      1      26.0      288      1      3
2      0      152.0      254      0      1
3      1      156.0      76      0      2
4      0      152.0      294      0      2
...    ...    ...    ...    ...    ...
11504793      1      26.0      218      0      3
11504794      1      152.0      115      1      1
11504795      0      152.0      189      0      1
11504796      1      26.0      274      1      4
11504797      0      152.0      189      0      1

      Annual_Premium_Binned
0      6
1      5
2      3
3      0
4      3
...    ...
11504793      2
11504794      2

```



```

11504795          0
11504796          4
11504797          3

```

```
[11504798 rows x 11 columns]
```

```
[43]: df_new.isnull().sum()
```

```

[43]: id          0
      Gender      0
      Driving_License  0
      Previously_Insured  0
      Vehicle_Age  0
      Vehicle_Damage  0
      Policy_Sales_Channel  0
      Vintage      0
      Response     0
      Age_Binned   0
      Annual_Premium_Binned  2307
      dtype: int64

```

```
[44]: df_new.dtypes
```

```

[44]: id          int64
      Gender      int64
      Driving_License  int64
      Previously_Insured  int64
      Vehicle_Age      int64
      Vehicle_Damage    int64
      Policy_Sales_Channel  float64
      Vintage          int64
      Response          int64
      Age_Binned        int64
      Annual_Premium_Binned  category
      dtype: object

```

1 Handling Missing Values in ‘Annual_Premium_Binned’

To handle missing values in the ‘Annual_Premium_Binned’ column of `df_new`, we first determine the mode (most frequent value) of the column using `df_new['Annual_Premium_Binned'].mode()[0]`. Next, we fill the missing values in the column with this mode value using `fillna(mode_value, inplace=True)`. Finally, to ensure the column contains integer values, we convert it to integer type using `astype(int)`.

```

[45]: mode_value = df_new['Annual_Premium_Binned'].mode()[0]
      df_new['Annual_Premium_Binned'].fillna(mode_value, inplace=True)

```

```
# Ensure the column is integer type
df_new['Annual_Premium_Binned'] = df_new['Annual_Premium_Binned'].astype(int)
```

/tmp/ipykernel_13/2197081812.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df_new['Annual_Premium_Binned'].fillna(mode_value, inplace=True)
```

```
[46]: pip install xgboost
```

Requirement already satisfied: xgboost in /usr/local/lib/python3.10/site-packages (2.1.0)

Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.10/site-packages (from xgboost) (2.20.5)

Requirement already satisfied: numpy in /usr/local/lib/python3.10/site-packages (from xgboost) (1.26.4)

Requirement already satisfied: scipy in /usr/local/lib/python3.10/site-packages (from xgboost) (1.13.1)

WARNING: Running pip as the 'root' user can result in broken permissions

and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: <https://pip.pypa.io/warnings/venv>

[notice] A new release of pip is available: 23.0.1 -> 24.1.2

[notice] To update, run:

```
pip install --upgrade pip
```

Note: you may need to restart the kernel to use updated packages.

2 XGBoost Workflow Overview

XGBoost, or eXtreme Gradient Boosting, sequentially enhances weak models like decision trees to minimize errors and improve predictions. By combining their outputs, it creates a powerful ensemble model that excels in accuracy and efficiency, making it widely used across diverse fields for its superior performance in complex data analysis tasks.

2.0.1 Steps:

Parameter Distributions:

- Define `param_dist` to specify the range of values for hyperparameters (`n_estimators`, `max_depth`, `learning_rate`, `gamma`, `reg_alpha`, `reg_lambda`) used in `RandomizedSearchCV`.

XGBoost Model Initialization:

- Initialize an `XGBClassifier` with a set `random_state` for reproducibility.

RandomizedSearchCV Setup:

- Configure `RandomizedSearchCV` to perform hyperparameter tuning with `n_iter=50` iterations, 5-fold cross-validation (`cv=5`), and using accuracy (`scoring='accuracy'`) as the evaluation metric.

Model Fitting:

- Fit `random_search` to the training data (`X_train`, `y_train`) to explore different hyperparameter combinations and identify the best-performing model.

Best Model Evaluation:

- Retrieve the best parameters (`best_params`) and best model (`best_model`) identified by `RandomizedSearchCV`.
- Make predictions on the test set (`X_test`) using `best_model`.
- Evaluate the model's performance using accuracy score, confusion matrix, and classification report.

This approach optimizes the XGBoost model's hyperparameters to achieve higher accuracy and robustness on unseen data, leveraging randomized search and cross-validation for effective model tuning.

```
[48]: from sklearn.model_selection import RandomizedSearchCV
      from scipy.stats import randint, uniform

      # Define the parameter distributions for RandomizedSearchCV
      param_dist = {
          'n_estimators': randint(50, 200), # Random integer values between 50 and
          ↪200
          'max_depth': randint(3, 8), # Random integer values between 3 and 7
          'learning_rate': uniform(0.001, 0.1), # Random float values between 0.001
          ↪and 0.1
          'gamma': uniform(0, 0.5), # Random float values between 0 and 0.5
          'reg_alpha': uniform(0, 0.5), # Random float values between 0 and 0.5
          'reg_lambda': uniform(1, 2), # Random float values between 1 and 2
      }

      # Initialize the XGBoost model
      xgb_model = XGBClassifier(random_state=42)
```

```

# Perform RandomizedSearchCV with cross-validation
random_search = RandomizedSearchCV(estimator=xgb_model,
    ↪param_distributions=param_dist,
                                n_iter=50, # Number of parameter settings
    ↪that are sampled
                                cv=5, # 5-fold cross-validation
                                scoring='accuracy', # Use accuracy for
    ↪scoring
                                random_state=42, # Random state for
    ↪reproducibility
                                n_jobs=-1, # Use all available cores
                                verbose=1) # Print detailed messages

# Fit the RandomizedSearchCV object to the training data
random_search.fit(X_train, y_train)

# Get the best parameters and model
best_params = random_search.best_params_
best_model = random_search.best_estimator_

# Print the best parameters found
print("Best Parameters:", best_params)

# Make predictions with the best model
y_pred = best_model.predict(X_test)

# Evaluate the best model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

# Print evaluation results
print(f'Accuracy: {accuracy:.2f}')
print('Confusion Matrix:')
print(conf_matrix)
print('Classification Report:')
print(class_report)

```

Fitting 5 folds for each of 50 candidates, totalling 250 fits

```

/usr/local/lib/python3.10/site-
packages/joblib/externals/loky/process_executor.py:752: UserWarning: A worker
stopped while some jobs were given to the executor. This can be caused by a too
short worker timeout or by a memory leak.
    warnings.warn(

```

Best Parameters: {'gamma': 0.28163778598819184, 'learning_rate':

```
0.07055160864261276, 'max_depth': 7, 'n_estimators': 192, 'reg_alpha':
0.37777556927152434, 'reg_lambda': 1.457596330983245}
```

Accuracy: 0.88

Confusion Matrix:

```
[[2007659  10117]
 [ 268842  14342]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.99	0.94	2017776
1	0.59	0.05	0.09	283184
accuracy			0.88	2300960
macro avg	0.73	0.52	0.51	2300960
weighted avg	0.85	0.88	0.83	2300960

2.1 Preprocessing the Test Data

(Following the same what we did in case of Train data)

```
[49]: data = pd.read_csv("/kaggle/input/playground-series-s4e7/test.csv")
data
```

```
[49]:
```

	id	Gender	Age	Driving_License	Region_Code	\
0	11504798	Female	20	1	47.0	
1	11504799	Male	47	1	28.0	
2	11504800	Male	47	1	43.0	
3	11504801	Female	22	1	47.0	
4	11504802	Male	51	1	19.0	
...	
7669861	19174659	Male	57	1	28.0	
7669862	19174660	Male	28	1	50.0	
7669863	19174661	Male	47	1	33.0	
7669864	19174662	Male	30	1	28.0	
7669865	19174663	Male	23	1	46.0	

	Previously_Insured	Vehicle_Age	Vehicle_Damage	Annual_Premium	\
0	0	< 1 Year	No	2630.0	
1	0	1-2 Year	Yes	37483.0	
2	0	1-2 Year	Yes	2630.0	
3	1	< 1 Year	No	24502.0	
4	0	1-2 Year	No	34115.0	
...	
7669861	0	1-2 Year	Yes	51661.0	
7669862	1	< 1 Year	No	25651.0	
7669863	1	1-2 Year	No	2630.0	
7669864	0	< 1 Year	Yes	38866.0	

7669865	1	< 1 Year	No	27498.0
---------	---	----------	----	---------

	Policy_Sales_Channel	Vintage
0	160.0	228
1	124.0	123
2	26.0	271
3	152.0	115
4	124.0	148
...
7669861	124.0	109
7669862	152.0	184
7669863	138.0	63
7669864	124.0	119
7669865	152.0	79

[7669866 rows x 11 columns]

```
[50]: columns_to_drop = ['Region_Code']

# Using the `drop` method
data_new = data.drop(columns=columns_to_drop)
```

```
[51]: age_bins = [10, 20, 30, 40, 50, 60, 70, 80, 90]
#premium_bins = [0,10000, 20000, 30000, 40000, 50000, 60000, 70000, 80000,
↳90000,100000]

# Apply binning using cut
data_new['Age_Binned'] = pd.cut(data_new['Age'], bins=age_bins,
↳labels=[f'{age}-{age+10}' for age in age_bins[:-1]])

data_new.drop(columns=['Age'], inplace=True)
```

```
[52]: # Assuming df is your DataFrame
bins = [0, 10000, 20000, 30000, 40000, 50000, 60000, 70000, 80000, 90000,
↳100000]
labels = range(len(bins)-1)

data_new['Annual_Premium_Binned'] = pd.cut(data_new['Annual_Premium'],
↳bins=bins, labels=labels)
```

```
[53]: import pandas as pd
from sklearn.preprocessing import LabelEncoder

# Assuming df is your DataFrame
label_encoder = LabelEncoder()

# Columns to encode
```

```
columns_to_encode = ['Gender', 'Vehicle_Age', 'Vehicle_Damage', 'Age_Binned']

# Apply label encoding to each column
for column in columns_to_encode:
    data_new[column] = label_encoder.fit_transform(data_new[column])
```

```
[54]: data_new.drop(columns=['Annual_Premium'], inplace=True)
data_new
```

```
[54]:
```

	id	Gender	Driving_License	Previously_Insured	Vehicle_Age	\
0	11504798	0	1	0	1	
1	11504799	1	1	0	0	
2	11504800	1	1	0	0	
3	11504801	0	1	1	1	
4	11504802	1	1	0	0	
...	
7669861	19174659	1	1	0	0	
7669862	19174660	1	1	1	1	
7669863	19174661	1	1	1	0	
7669864	19174662	1	1	0	1	
7669865	19174663	1	1	1	1	

	Vehicle_Damage	Policy_Sales_Channel	Vintage	Age_Binned	\
0	0	160.0	228	0	
1	1	124.0	123	3	
2	1	26.0	271	3	
3	0	152.0	115	1	
4	0	124.0	148	4	
...	
7669861	1	124.0	109	4	
7669862	0	152.0	184	1	
7669863	0	138.0	63	3	
7669864	1	124.0	119	1	
7669865	0	152.0	79	1	

	Annual_Premium_Binned
0	0
1	3
2	0
3	2
4	3
...	...
7669861	5
7669862	2
7669863	0
7669864	3
7669865	2

[7669866 rows x 10 columns]

```
[55]: mode_value = data_new['Annual_Premium_Binned'].mode()[0]
data_new['Annual_Premium_Binned'].fillna(mode_value, inplace=True)

# Ensure the column is integer type
data_new['Annual_Premium_Binned'] = data_new['Annual_Premium_Binned'].
    ↪astype(int)
```

/tmp/ipykernel_13/2780489815.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
data_new['Annual_Premium_Binned'].fillna(mode_value, inplace=True)
```

2.2 Prediction

```
[56]: import pandas as pd
from xgboost import XGBClassifier

# Assuming 'data_new' is your new dataframe for prediction
# data_new = pd.read_csv('path_to_new_data.csv') # Load your new data if not
    ↪already loaded

# Predict probabilities for data_new using the best model
y_pred_proba = best_model.predict_proba(data_new)[:, 1] # Predict probability
    ↪of positive response

# Add 'Response' column with predicted probabilities to data_new
data_new['Response'] = y_pred_proba

# Save to CSV file with only 'id' and 'Response' columns
output_file = 'predictions.csv'
data_new[['id', 'Response']].to_csv(output_file, index=False)

print(f"Predictions saved to '{output_file}'")
```

Predictions saved to 'predictions.csv'


```
[57]: prediction = pd.read_csv("predictions.csv")
      prediction.sample(n=10)
```

```
[57]:
```

	id	Response
6851316	18356114	0.049780
1231893	12736691	0.499255
4835523	16340321	0.244009
6284369	17789167	0.000577
1462131	12966929	0.212542
6512807	18017605	0.002250
1348706	12853504	0.215294
126158	11630956	0.274941
5066806	16571604	0.123153
2860125	14364923	0.000431